

Parallel and Distributed Systems - Appunti

Francesco Lorenzoni

September 2024

Contents

I Introduction to SDC	5
1 Basic Concepts	9
1.1 Parallel Computing	9
1.1.1 Current usages	9
2 Compilation - Leiserson MIT	11
2.1 Interpreters vs Compilers	11
2.2 Cache	11
2.3 Compiler Optimization	12
2.4 Parallelizing	12
2.5 Tiling	14
2.6 Where have we gotten so far - Further optimizations	15
2.6.1 Recursion in Tiling	15
2.6.2 Vectorization Flags	15
3 Parallel Architectures	17
3.1 Flynn's Taxonomy	17
3.1.1 MIMD Architectures	17
3.2 Classifying on cores count	18
3.3 Programming Parallel Architectures	18
3.3.1 Shared-Memory	18
3.3.2 Distributed-Memory	18
3.3.3 Summary	18
3.4 Suggested Readings	19
4 Shared Memory Architectures	21
4.1 Von Neumann Bottleneck	21
4.1.1 Caches	21
4.2 Locality Principle	22
4.2.1 Measuring CPU time with caches	22
4.2.2 Cache Algorithms	22
4.2.3 Cache mapping and eviction strategies	22
4.2.4 Cache Write Policies	23
4.2.5 Cache Coherence	23
4.3 Suggested readings	23
4.4 Advanced Processors and Technologies	23
4.4.1 Superscalar CPUs	23
4.4.2 HW Multithreading	23
4.5 Programming Shared Memory Systems	24
4.5.1 Threads are the way to go	24
4.5.2 Data-race	24
4.5.3 False Sharing	24
4.5.3.1 Padding	25
4.5.3.2 Local variables	25
4.6 SIMD and Vectorization	26
4.6.1 AVX registers	26
4.7 Vectorizing code	26
4.7.1 AoS vs SoA	26
4.7.2 Loops - Compiler Optimization	27

5	Distributed Memory Systems	29
5.1	Interconnection Networks	29
5.1.1	Terminology	29
5.1.2	Examples	31
5.1.2.1	Hypercube	31
5.1.3	Criteria for Evaluation - Summary	31
5.1.3.1	Simple Communication cost model	31
5.1.4	More advanced topologies	31
5.1.4.1	Fat-tree	32
5.1.4.2	Dragonfly	32
5.2	Computation-to-Communication Overlap	32
5.2.1	Synchronous vs Asynchronous Communication	32
5.3	Foster's Parallel Algorithm Design Method	33

Part I

Introduction to SDC

1 Basic Concepts	9
1.1 Parallel Computing	9
1.1.1 Current usages	9
2 Compilation - Leiserson MIT	11
2.1 Interpreters vs Compilers	11
2.2 Cache	11
2.3 Compiler Optimization	12
2.4 Parallelizing	12
2.5 Tiling	14
2.6 Where have we gotten so far - Further optimizations	15
2.6.1 Recursion in Tiling	15
2.6.2 Vectorization Flags	15
3 Parallel Architectures	17
3.1 Flynn's Taxonomy	17
3.1.1 MIMD Architectures	17
3.2 Classifying on cores count	18
3.3 Programming Parallel Architectures	18
3.3.1 Shared-Memory	18
3.3.2 Distributed-Memory	18
3.3.3 Summary	18
3.4 Suggested Readings	19
4 Shared Memory Architectures	21
4.1 Von Neumann Bottleneck	21
4.1.1 Caches	21
4.2 Locality Principle	22
4.2.1 Measuring CPU time with caches	22
4.2.2 Cache Algorithms	22
4.2.3 Cache mapping and eviction strategies	22
4.2.4 Cache Write Policies	23
4.2.5 Cache Coherence	23
4.3 Suggested readings	23
4.4 Advanced Processors and Technologies	23
4.4.1 Superscalar CPUs	23
4.4.2 HW Multithreading	23
4.5 Programming Shared Memory Systems	24
4.5.1 Threads are the way to go	24
4.5.2 Data-race	24
4.5.3 False Sharing	24
4.6 SIMD and Vectorization	26
4.6.1 AVX registers	26
4.7 Vectorizing code	26
4.7.1 AoS vs SoA	26
4.7.2 Loops - Compiler Optimization	27
5 Distributed Memory Systems	29
5.1 Interconnection Networks	29
5.1.1 Terminology	29
5.1.2 Examples	31
5.1.3 Criteria for Evaluation - Summary	31
5.1.4 More advanced topologies	31
5.2 Computation-to-Communication Overlap	32
5.2.1 Synchronous vs Asynchronous Communication	32
5.3 Foster's Parallel Algorithm Design Method	33

Chapter 1

Basic Concepts

Fun fact: SPM stands for *Software Paradigms and Models*, the historical name of the course

1.1 Parallel Computing

Definition 1.1 (Parallel Computing) *the practice of using multiple processors in parallel to solve problems more quickly than with a single processor. It implies the capability of:*

- ◊ identifying and exposing parallelism in algorithms and software systems
- ◊ understanding the costs, benefits, and limitations of a given parallel implementation

1.1.1 Current usages

The motivation for parallel computing is the need to solve larger and more complex problems in less time, typically *simulation* ones, but not only. Besides, today, even from the single machine perspective, there exists no more the single processor architecture, so parallel addresses also exploiting the multiple cores available in a single machine.

- ◊ Big Data Analytics (BDAs)
- ◊ HPC and/or AI

Besides also the *Moore's law* indicates another motivation:

Definition 1.2 *Gordon Moore, co-founder of Intel, observed that the number of transistors on a chip doubles every 18-24 months, leading to a doubling of the performance of the chip.*

However, even if the number of transistors on a chip continues to increase, we started to face the problem of powering simultaneously all the transistors, leading to the *power wall* problem. It was estimated in the early 00s that the *power density* of a chip would reach the power density of a nuclear reactor by 2020, and then the power density of the sun in a while. This was the main reason for the shift from single-core to **multi-core** chips (**CMPs**).

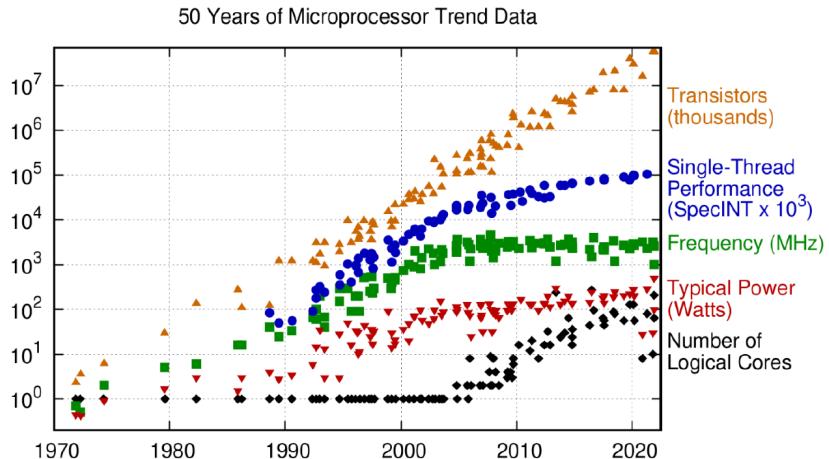


Figure 1.1: Microprocessors in the last 30 years

Single thread performance is increasing slowly, while the Frequency is stable. Moore's law is still valid if we account the number of cores.

Multicore processors help reducing power for this reason:

1. Doubling the number of cores doubles the performance, but also power ☺
2. Doubling the number of cores and *halving* Voltage and Frequency, leaves the same performance unaltered, but the power consumption is reduced by a factor of 4. ☺

To fully exploit the potential of multicore processors, programmers need to **parallelize** our software.

There also forms of parallelization under-the-hood, which make the parallelization transparent to the developer. There also libraries that help in parallelizing the code, such as *OpenMP* or *FastFlow*.

There also Heterogeneous CMPs which integrate different processor cores in a single chip, but they are more complex to handle. Common examples are the integration of a GPU in the chip, or the integration of a *big.LITTLE* architecture, which integrates high-performance cores with low-power cores. Real-world uses are some ARM processors, or the Apple M1.

Chapter 2

Compilation - Leiserson MIT

2.1 Interpreters vs Compilers

Interpreted languages are more versatile, but much slower.

```
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        for (int k = 0; k < n; k++) {
            C[i][j] += A[i][k] +B[k][j];
        }
    }
}
```

This code executed using Clang/LLVM 5.0 takes 1156s (19m 16s) to execute, about **2x** times faster than Java and **18x** times than python

2.2 Cache

```
for (int i = 0; i < n; i++) {
    for (int k = 0; k < n; k++) {
        for (int j = 0; j < n; j++) {
            C[i][j] += A[i][k] +B[k][j];
        }
    }
}
```

Loop order (outer to inner)	Running time (s)	Last-level-cache miss rate
i, j, k	1155.77	7.7%
i, k, j	177.68	1.0%
j, i, k	1080.61	8.6%
j, k, i	3056.63	15.4%
k, i, j	179.21	1.0%
k, j, i	3032.82	15.4%

We can change the order of the loops without changing the result, but the performance can change.

Figure 2.1: Performance against loop order

As you can see, there is a huge difference in the running time of the loop depending on the loops ordering. This is due to **caching**, which consists in storing in a fast-access memory previously accessed memory lines.



Figure 2.1: Memory layout for matrix rows

Matrices are stored in memory in row-major order, so the first loop should iterate over the rows of the matrix, to exploit the cache.

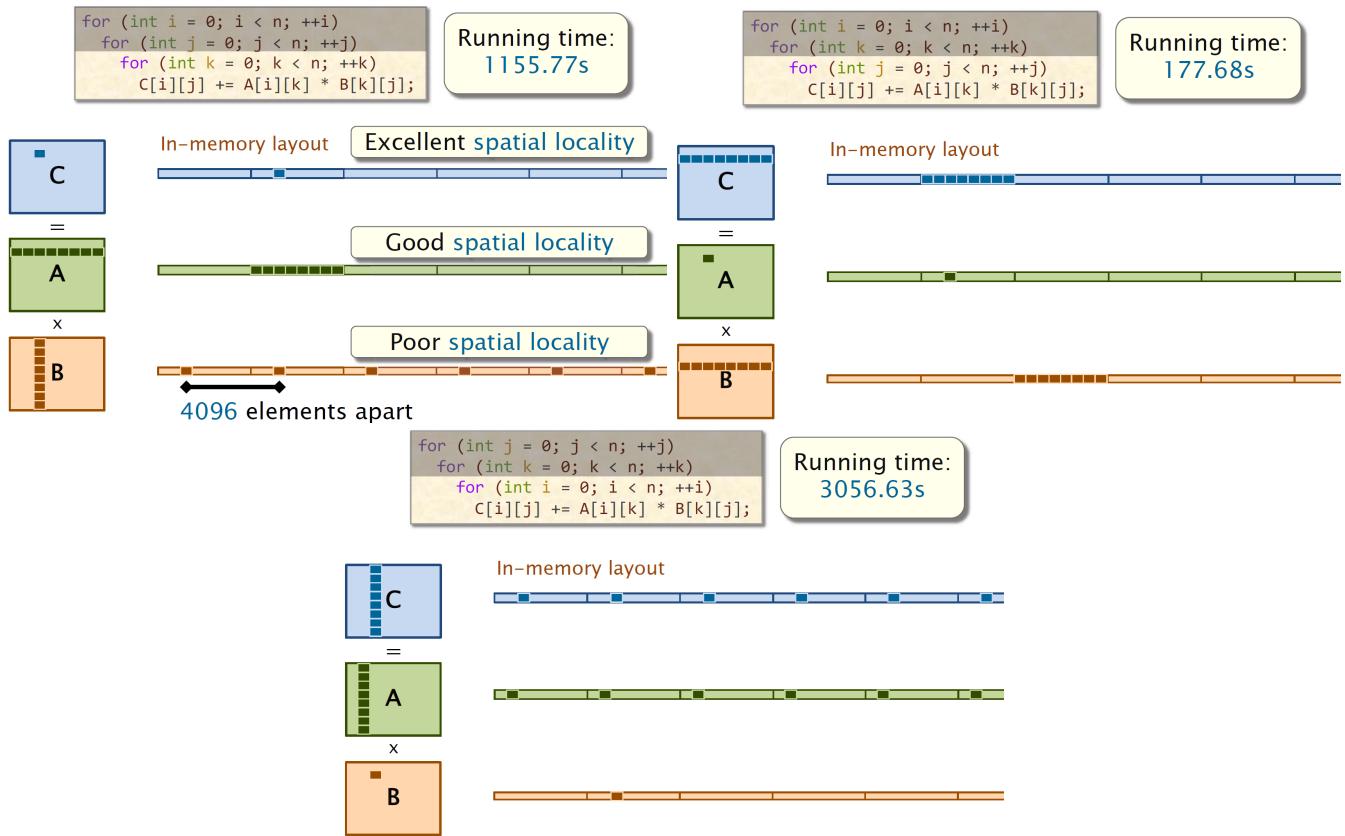


Figure 2.2: Memory layout and spaciality implications

2.3 Compiler Optimization

Clang offers a lot of optimization flags, like `-O3` which enables all the optimizations. The compiler can also unroll loops, which means that it can execute multiple iterations of the loop in parallel. This can be done only if the number of iterations is known at compile time. There are also `-Os` which optimizes for size, and `-Og` which generates debug information. There's plenty of them, for various uses.

Opt. level	Meaning	Time (s)
<code>-O0</code>	Do not optimize	177.54
<code>-O1</code>	Optimize	66.24
<code>-O2</code>	Optimize even more	54.63
<code>-O3</code>	Optimize yet more	55.58

Figure 2.3: Optimization flags and relative performance

2.4 Parallelizing

Even after all these tweaks, we are still using only one of the 9 cores of the CPU. So...

```
cilk_for (int i = 0; i < n; i++) {
    for (int k = 0; k < n; k++) {
        cilk_for (int j = 0; j < n; j++) {
            C[i][j] += A[i][k] + B[k][j];
        }
    }
}
```

We don't have to know what's behind the `cilk_for` keyword, but it will parallelize the `for` loop execution.

But which for loops should we parallelize?

Parallelizing all three would cause multiple threads to access the same memory, which would be messy.

A **rule of thumb** is to parallelize the **outermost** loop, which is the one that iterates over the rows of the matrix.

This is demonstrated by the following slide.

Parallel i loop

```
cilk_for (int i = 0; i < n; ++i)
    for (int k = 0; k < n; ++k)
        for (int j = 0; j < n; ++j)
            C[i][j] += A[i][k] * B[k][j];
```

Running time: 3.18s

Parallel j loop

```
for (int i = 0; i < n; ++i)
    for (int k = 0; k < n; ++k)
        cilk_for (int j = 0; j < n; ++j)
            C[i][j] += A[i][k] * B[k][j];
```

Running time: 531.71s

Parallel i and j

```
cilk_for (int i = 0; i < n; ++i)
    for (int k = 0; k < n; ++k)
        cilk_for (int j = 0; j < n; ++j)
            C[i][j] += A[i][k] * B[k][j];
```

Running time: 10.64s

Rule of Thumb
Parallelize outer
loops rather than
inner loops.

Figure 2.3: Parallelizing only the outermost loop leads to optimal performance

2.5 Tiling

Well, the possible optimizations ain't over ☺. Consider the first picture and let's dig into some math. How many

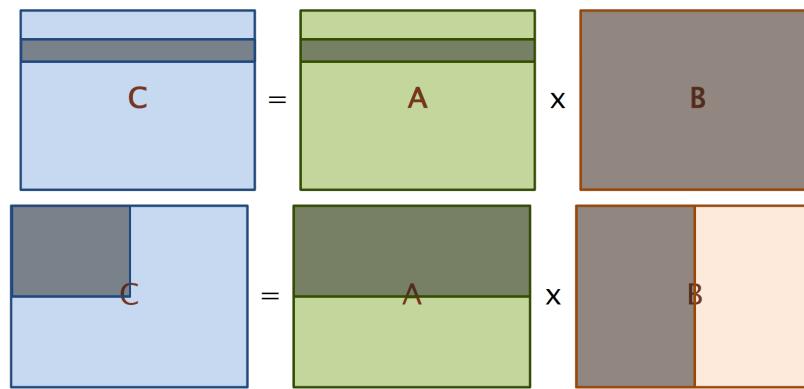


Figure 2.4: Tiling

memory accesses must the looping perform to fully compute 1 row of C?

$$4096 \cdot 1 = 4096 \text{ writes to } C \quad (2.1)$$

$$4096 \cdot 1 = 4096 \text{ reads from } C \quad (2.2)$$

$$4096 \cdot 4096 = 1.6777216 \cdot 10^6 \text{ reads from } B \quad (2.3)$$

$$1.6777216 + 4096 + 4096 = 1.6785408 \cdot 10^6 \text{ total memory accesses} \quad (2.4)$$

But if we consider instead computing a 64×64 block of C we can shrink down the number of memory accesses to half a million:

$$64 \cdot 64 = 4096 \text{ writes to } C \quad (2.5)$$

$$64 \cdot 4096 = 262144 \text{ reads from } A \quad (2.6)$$

$$4096 \cdot 64 = 262144 \text{ reads from } B \quad (2.7)$$

$$262144 + 262144 + 4096 = 528384 \text{ total memory accesses} \quad (2.8)$$

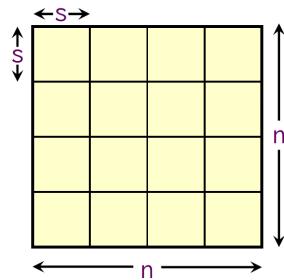
$$(2.9)$$

But in general, which would the optimal block size be? The only way is to experiment.

Why?

`cilk_for (int ih = 0; ih < n; ih += s)
 cilk_for (int jh = 0; jh < n; jh += s)
 for (int kh = 0; kh < n; kh += s)
 for (int il = 0; il < s; ++il)
 for (int kl = 0; kl < s; ++kl)
 for (int jl = 0; jl < s; ++jl)
 C[ih+il][jh+jl] += A[ih+il][kh+kl] * B[kh+kl][jh+jl];`

Tuning parameter
How do we find the right value of s ? Experiment!



Tile size	Running time (s)
4	6.74
8	2.76
16	2.49
32	1.74
64	2.33
128	2.13

Figure 2.5: Tile size

2.6 Where have we gotten so far - Further optimizations

Version	Implementation	Running time (s)	Relative speedup	Absolute Speedup	GFLOPS	Percent of peak
1	Python	21041.67	1.00	1	0.006	0.001
2	Java	2387.32	8.81	9	0.058	0.007
3	C	1155.77	2.07	18	0.118	0.014
4	+ interchange loops	177.68	6.50	118	0.774	0.093
5	+ optimization flags	54.63	3.25	385	2.516	0.301
6	Parallel loops	3.04	17.97	6,921	45.211	5.408
7	+ tiling	1.79	1.70	11,772	76.782	9.184

Implementation	Cache references (millions)	L1-d cache misses (millions)	Last-level cache misses (millions)
Parallel loops	104,090	17,220	8,600
+ tiling	64,690	11,777	416

The tiled implementation performs about **62%** fewer cache references and incurs **68%** fewer cache misses.

Figure 2.6: Comparison of the various optimizations

2.6.1 Recursion in Tiling

Tiling may be also implemented as a divide-and-conquer algorithm exploiting recursion. This yields slightly better performance, but requires to tune the recursion base case **threshold**. Having a too small threshold would lead to a lot of overhead, due to many function invocations.

2.6.2 Vectorization Flags

There may be also flags to enable instructions specific of a given architecture:

- ◊ **-mavx**: Use Intel AVX vector instructions.
- ◊ **-mavx2**: Use Intel AVX2 vector instructions.
- ◊ **-mfma**: Use fused multiply-add vector instructions.
- ◊ **-march=<string>**: Use whatever instructions are available on the specified architecture.
- ◊ **-march=native**: Use whatever instructions are available on the architecture of the machine doing compilation.

Due to restrictions on floating-point arithmetic, additional flags, such as **-ffast-math**, might be needed for these vectorization flags to have an effect

You could also use AVX Intrinsic Instructions that provide access to hardware vector operations. They are available in C and C++. software.intel.com/sites/landingpage/IntrinsicsGuide.

These may help further more, but we are getting very closer to the hardware.

Chapter 3

Parallel Architectures

3.1 Flynn's Taxonomy

There are various classifications possible for parallel architectures, but the most common one is the one based on the **Flynn's Taxonomy**.

This taxonomy is based on the number of **instructions** and **data streams**

- ◊ SISD (Single Instruction, Single Data): the classic Von Neumann architecture, with a single processor executing a single instruction on a single data stream.
- ◊ SIMD (Single Instruction, Multiple Data): a single instruction is broadcasted to multiple processors, each of which operates on a different data stream. This is the architecture of GPUs.
- ◊ MISD (Multiple Instruction, Single Data): multiple processors execute different instructions on the same data stream. This is not common in practice.
- ◊ MIMD (Multiple Instruction, Multiple Data): multiple processors execute different instructions on different data streams. This is the most common architecture for parallel systems.

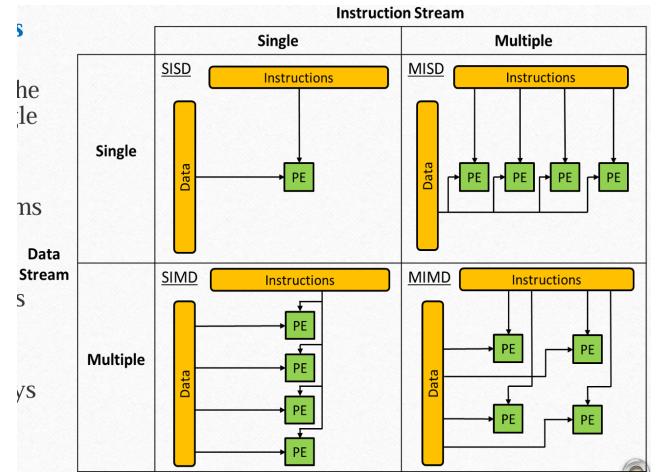


Figure 3.1: Flynn taxonomy

3.1.1 MIMD Architectures

A set of PEs (Processing Elements) simultaneously execute different instructions on different data streams. Each processor can execute all instructions. This is the most common architecture for parallel systems.

This architecture can be further classified considering memory organization and interconnection (between PE and MM) topologies.

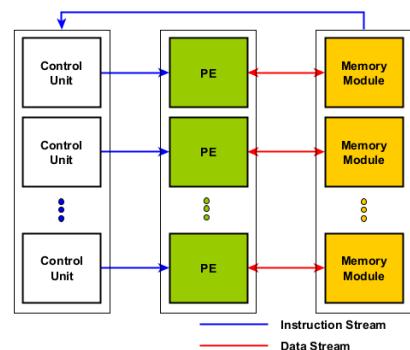


Figure 3.2: MIMD architecture

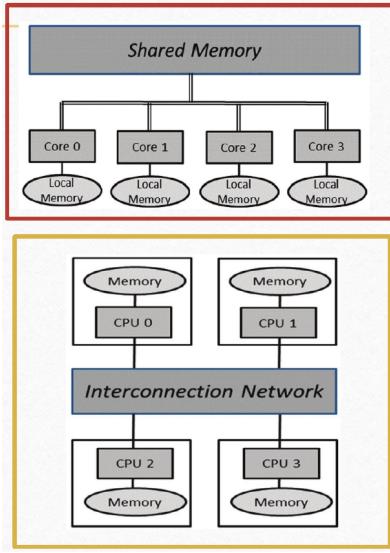


Figure 3.1: Classification based on memory

◊ **Shared Memory MIMD**: all processors share the same memory space, and can access it directly. This is the most common architecture for multi-core processors (“*multiprocessors*”).

- **UMA (Uniform Memory Access)**: all processors access memory with the same latency. SMP - Symmetric Multiprocessor

- **NUMA (Non-Uniform Memory Access)**: processors access memory with different latencies.

Distributed Memory architectures are inherently NUMA

◊ **Distributed Memory MIMD**: each processor has its own memory space, and communicates with other processors through messages. This is the most common architecture for clusters.

Interconnection may be based on Ethernet, InfiniBand, etc.

Historically called “*multicomputers*”

3.2 Classifying on cores count

- ◊ $\mathcal{O}(10^1 \div 10^2)$ cores, for a single multiprocessor chip (CMP)
- ◊ $\mathcal{O}(10^2 \div 10^3)$ cores, for a Shared Memory tightly-coupled multiprocessor
- ◊ $\mathcal{O}(10^3 \div 10^5)$ cores, for a Distributed Memory loosely-coupled multiprocessor (small to large compute clusters)
- ◊ $\mathcal{O}(10^5 \div 10^6)$ cores, top supercomputers (e.g. *Leonardo @cineca*)

3.3 Programming Parallel Architectures

3.3.1 Shared-Memory

In *Shared-Memory* systems the emphasis is on the memory organization, meaning the memory hierarchy and processor-memory interconnections, aiming to reduce the “*von Neumann bottleneck*”. Critical aspects are cache coherence, memory consistency and thread synchronization.

Programming models are based on *thread-level parallelism* to exploit the physical shared memory by means of **Shared Variables** programming models (e.g. OpenMP, Pthreads, Cilk).

Caches shared among cores are used to reduce the latency of memory accesses, and to exploit the *spatial locality* of data accesses, however they introduce the problem of *cache coherence*, with respective issue (more on this later).

3.3.2 Distributed-Memory

In *Distributed-Memory* systems the emphasis is on the interconnection network, aiming to reduce the “*communication bottleneck*”, so reducing latency and increasing bandwidth. Critical aspects are messaging protocols/libraries, routing.

Here the only parallelism exploitable is *process-level parallelism*, so the programming models are based on **message-passing** (e.g. MPI, POSIX socket, PVM).

Nowdays, each node is a CMP. Depending on the network and some other aspects, we may further classify these systems in Clusters, Cloud, geographical distributed systems, etc.

We are interested in systems with high performance network topologies and homogeneous nodes, like compute clusters.

A common example of application is the **Stencil Computation**, which is a common pattern in scientific computing, where each element of a matrix is computed as a function of its neighbors.

3.3.3 Summary

Distributed Memory systems are more scalable, costly and less energy efficient.

From the programming perspective, Shared Memory systems are easier to program and the physical shared memory can be used for fast communication between threads. However, locking and synchronization are critical points deserving

attention.

For what concerns Distributed Memory systems, the most important aspect is to reduce as much as possible the cost of communication (i.e. I/O), for instance overlapping computation and communication, reducing memory copies for I/O, using fast messaging protocols (e.g. RDMA).

3.4 Suggested Readings

- ◊ Chapter 1 - Section 1.2 - “Parallelism Basics” of Parallel Programming Concept and Practice book
- ◊ Chapter 3 - Section 3.2 - “Flynn’s Taxonomy” of Parallel Programming Concept and Practice book

Chapter 4

Shared Memory Architectures

Shared Memory Architectures are a type of MIMD architecture where all processors share the same memory space, and can access it directly. This is the most common architecture for multi-core processors (“*multiprocessors*”).

They mirror the Von Neumann architecture, with multiple processors sharing the same memory space.

4.1 Von Neumann Bottleneck

Definition 4.1 (von Neumann Bottleneck) *The von Neumann bottleneck is a limitation on throughput caused by the standard personal computer architecture. The term is named for John von Neumann, who is credited with developing the von Neumann architecture, in which programs and data are stored in the same memory. The bottleneck refers to the limited data transfer rate between a computer’s CPU and memory compared to the amount of memory.*

- Let's consider the Dot Product kernel (see code snippet)
 - If $n = 2^{30} \rightarrow 2 * n$ floating point operations (i.e., 2 GFlop), and $2 * n * 8B = 16GB$ data transferred from memory, then:
- $$t_{comp} = \frac{2 \text{ GFlop}}{384 \text{ GFlop/s}} = 5.2ms, \quad t_{mem} = \frac{16 \text{ GB}}{51.2 \text{ GB/s}} = 312.5ms$$
- If we overlap computation and memory data transfer, a lower bound of the execution time is
$$t_{exec} \geq \max(t_{comp}, t_{mem}) = 312.5ms$$
 - Achievable performance: $\frac{2 \text{ GFlop}}{312.5 \text{ ms}} = 6.4 \text{ GFlop/s}$
 - i.e., less than 2% of peak compute performance
 - Considering our architecture, «Dot Product» is **memory bound**

```
// Dot Product
double dotp = 0.0;
for (int i = 0; i < n; i++)
    dotp += u[i] * v[i];
```

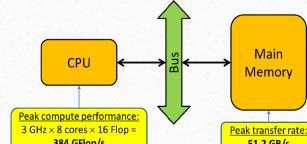


Figure 4.1: von Neumann Bottleneck example

In the example in Figure 4.1, the CPU is waiting for the data to be loaded from memory, which is a slow operation, leading to exploiting only the 2% of the CPU capabilities.

4.1.1 Caches

Back in the day, the solution was to “*move the data closer to the CPU*”, introducing **memory hierarchy** and **caches**.

Usually, L1 and L2 are private to each core, while L3 is shared among all cores.

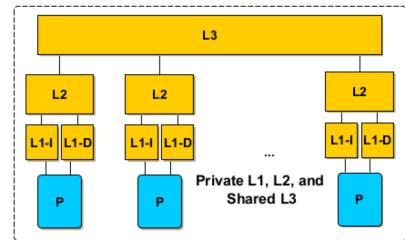


Figure 4.3: Caches hierarchy

If a matrix of the previous example in Fig. 4.1 is entirely stored in cache, then the achievable performance is 223 GFLOPS, which is about 60% of the peak power.

If the matrices do not fit in the cache, the performance drops since we need to trap to the main memory to fetch

missing data.

So, the problem shifts to understand how to exploit the cache to its fullest.

4.2 Locality Principle

The locality principle is the driving force that makes the cache work. Locality increases the probability of reusing data blocks that were previously moved from level n to level $n - 1$.

- ◊ **Temporal Locality:** if a data is accessed, it is likely to be accessed again soon.
Cache mapping strategy (Direct/associative) and the replacement policy (LRU, FIFO, Random etc.) are crucial to exploit temporal locality.
- ◊ **Spatial Locality:** if a data is accessed, it is likely that nearby data will be accessed soon.

4.2.1 Measuring CPU time with caches

$$CPU_{time} = ClockCycles \cdot ClockCycleTime = IC \cdot ClockCycleTime \quad (4.1)$$

- ◊ IC: Instruction Count (number of instructions executed)
 - $IC = IC_{CPU} + IC_{MEM}$
- ◊ CPI: Cycles Per Instruction
 - $CPI = \frac{ClockCycles}{IC}$
 - $CPI = \frac{IC_{CPU}}{IC} \cdot CPI_{CPU} + \frac{IC_{MEM}}{IC} \cdot CPI_{MEM}$ where CPI_{CPU} are the average cycles per ALU instruction and CPI_{MEM} are the average cycles per memory instruction.
 - Considering that each memory instruction may generate a cache hit or miss with a given probability, and naming *HitRate* the probability of a cache hit, we can write

$$CPI_{MEM} = HitRate \cdot CPI_{MEM-Hit} + (1 - HitRate) \cdot CPI_{MEM-Miss} \quad (4.2)$$

4.2.2 Cache Algorithms

1. *What do we load from main memory?*
2. *Where do we store it in the cache?*
3. *Cache is full, what should we evict?*

At the beginning of the second half of the 4th lecture, the professor displays how to VPN in the unipi and then ssh to the servers.

4.2.3 Cache mapping and eviction strategies

- ◊ **Direct-mapped** cache: each memory block can be placed in only one cache line.
- ◊ **n-way** set associative cache: each memory block can be placed in n cache lines.
- ◊ **Least Recently Used (LRU)** cache: the block that has been accessed the least recently is evicted.

Transposing Matrices

Suppose you have to multiply matrix A and B. If you access A as rows, you'll access B as columns.

Since matrices are stored in row-major order, you won't exploit spatial locality on B, and for each element of A you'll have a cache miss on B, creating the need to evict a line and load another one in cache.

Transposing B would solve the problem, since you would access B as rows, exploiting spatial locality.

Prof. Torquati displayed that transposing and then multiplying the matrices would lead to a 2x speedup.

4.2.4 Cache Write Policies

Data in cache may be inconsistent with the value in memory, leading to the need to write back the data to memory. There are two policies:

- ◊ **Write-through:** data is written to both cache and memory. It is simple but slow.
- ◊ **Write-back:** data is written only to the cache, and then to memory when the block is evicted. It is faster but more complex.
 - Caches mark data in the cache as dirty (Dirty bit)
 - When a dirty line is evicted, it is written in main memory
 - A store write buffer is generally used to reduce the cost of cache writes

4.2.5 Cache Coherence

With private caches per core, it is possible to have several copies of shared data in distinct caches, each cache stores a different value for a single address location.

Definition 4.2 (Cache inconsistency) *Two caches store different values for the same variable.*

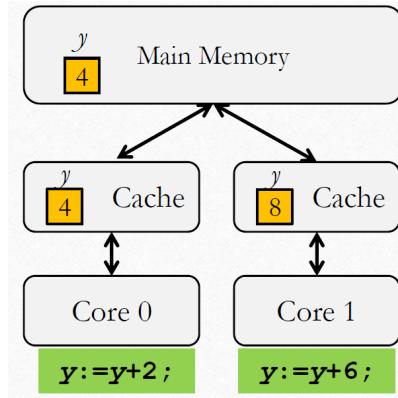


Figure 4.4: Cache inconsistency problem

Hardware firmware automatically solves this problem, but it is important to know that it exists. The algorithms responsible for this are called **cache coherence protocols**, a famous one is the **MESI** protocol, but we ain't going to study it. Note however that the cache coherence protocol granularity is the *cache line*, not the *single variable*.

4.3 Suggested readings

- ◊ Chapter 3.1 of the Parallel Programming Concept and Practice book
- ◊ Cache-coherence protocols: “A Primer on Memory Consistency and Cache Coherence” Daniel J. Sorin, Mark D. Hill, and David A. Wood

4.4 Advanced Processors and Technologies

4.4.1 Superscalar CPUs

Superscalar CPUs are designed to execute multiple instructions from a single process/thread simultaneously to improve performance and CPU utilization

The processor fetches multiple instructions concurrently in a single clock cycle and executes them out-of-order, exploiting instruction-level parallelism. Results are then re-ordered to ensure they are written back to the register file or memory in the correct program order

However, in sequential programs, the number of instructions that are independent are small thus, the exploited parallelism is low. To overcome this limitation, **SMT** (*Simultaneous Multi-Threading*) has been added in superscalar processors to execute multiple instructions from multiple threads of control simultaneously.

Hyperthreading is Intel's implementation of SMT

4.4.2 HW Multithreading

HW multithreading enables a single core to execute multiple threads concurrently. There are two main types of HW multithreading:

- ◊ **Fine-grained multithreading:** the processor switches between threads at each cycle (instruction level).

- ◊ **Coarse-grained multithreading:** the processor switches between threads only when the thread in execution causes a stall.

Each thread has its own set of registers and program counter; the processor maintains the context of each thread to quickly switch between them. However, they share the same cache and execution units. For the OS, each context is seen as a logical core.

4.5 Programming Shared Memory Systems

4.5.1 Threads are the way to go

Thread creation is more lightweight and faster (from 3x to 5x) than process creation¹, and threads share the same memory space, so they can communicate easily. Creating a thread takes $\mathcal{O}(10^4)$ cycles in C/C++.

4.5.2 Data-race

Definition 4.3 (Data Race) Scenario that occurs when two threads access a shared variable simultaneously and at least one of the accesses is a write, and the accesses are not guarded by a synchronization operation.

¹A `fork` system call requires copying (`memcpy`) more data, e.g. page table

DRs produce non-deterministic results, and are hard to debug, since they depend on the thread scheduling. To avoid this issue, we may use *mutexes*, *condition variables*, *semaphores*, *atomic operations*, etc.

We will discuss all four, except *semaphores*, which are more common in processes synchronization.

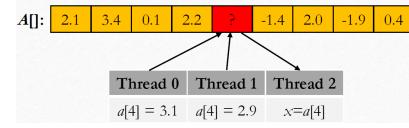


Figure 4.5: Data Race example

4.5.3 False Sharing

Caches are organized in cache lines, and cache coherence is managed at the cache line granularity.

Definition 4.4 (False Sharing) Scenario that occurs when two threads access different variables that reside on the same cache line. The cache coherence protocol will invalidate the cache line, so actually the two threads won't share anything.

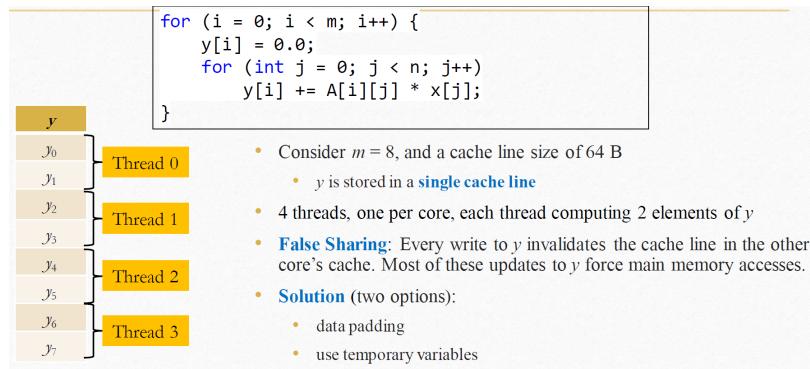


Figure 4.2: False Sharing slide

Different threads access different variables which however reside on the same cache line.

When False Sharing occurs, the performance of the system is degraded, since the cache line is invalidated and the cache coherence protocol keeps the cache line consistent among all copies.

Looking at Fig. 4.2, suppose that T_1 accesses y_2 : the entire y cache line is —marked as?— **invalidated**, so when T_1 accesses y_1 , it will have a cache miss, and the cache line will be **reloaded from memory** “moved from one core to another”²

²Prof. Torquati said this, not sure if saying it gets “reloaded from mem” is correct

4.5.3.1 Padding

```
# elapsed time (sequential_increment): 2.68141s
1073741824 1073741824
# elapsed time (false_sharing_increment_increment): 10.3552s
1073741824 1073741824
```

```
// 128-bit packed data that fits
// into a cache line of 64 B
// 2 threads, one modify ying
// the other thread modify yang
struct pack_t {
    uint64_t ying;
    uint64_t yang;
};
```

Using explicit data padding

```
// forcing ying and yang to lay in
// two distinct cache lines
struct pack_t {
    uint64_t ying;
    char padding[CACHELINE_SIZE_BYTE -
                 sizeof(uint64_t)];
    uint64_t yang;
};
```

Figure 4.3: False Sharing demo exec demo

Prof. Torquati displayed that the exec time of a dummy program which increments two distinct variables in a struct went from 2.6s to 10.3s, when going from sequential to parallel execution. The overhead is due to false sharing.

Fig. 4.3 also displays a possible solution to the problem: **padding** the struct with a dummy variable, so that the two variables are placed on different cache lines. Even though looks a bit hardcoded, it works indeed! Exec time, went from 2.6s to 2.9s, basically the same time, since the two variables are now on different cache lines; there is only some overhead due to the threads creation.

How can i understand if false sharing is happening in a complex program?

Test.

4.5.3.2 Local variables

Actually the better solution would be to use **local variables**, since they are stored in registers, and each thread has its own stack, so there is no sharing at all.

```
for (i = 0; i < m; i++) {
    float _y = 0.0;
    for (int j = 0; j < n; j++)
        _y += A[i][j] * x[j];
    y[i] = _y;
}
```

Figure 4.4: Local variable stored in registers

Compiler optimization

Modern compilers are able to optimize the code and avoid false sharing. If the example in Fig. 4.3 is compiled with `g++ -O3` (instead of `-O0`), the exec time goes back to normal.

However, it is not always possible to rely on the compiler, and in fact many times it does not work. So, it is better to pad the struct or use local variables, or avoid false sharing in general.

4.6 SIMD and Vectorization

SIMD - Single Instruction, Multiple Data. A parallel computing model that exploits data-level parallelism. There are two limitations which must be kept in mind:

1. ALUs are *limited in number*, so the number of operations that can be executed in parallel is limited.
2. All ALUs must execute the *same instruction*.

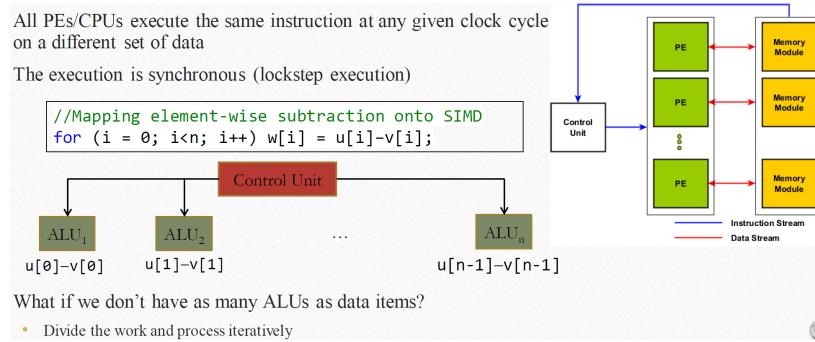


Figure 4.4: SIMD and ALUs

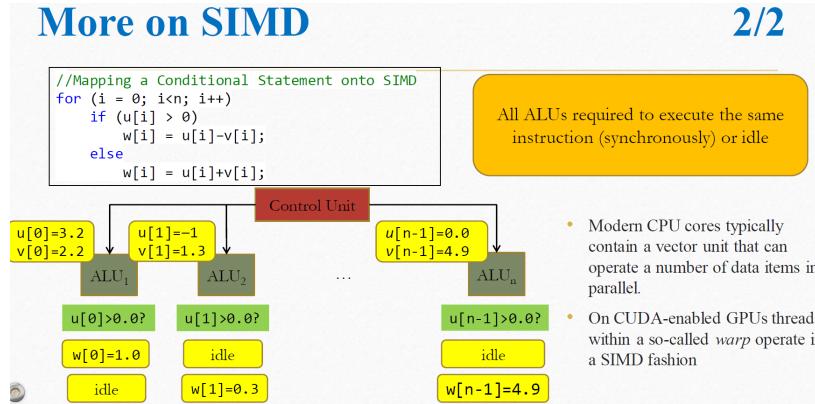


Figure 4.5: Here the ALUs either execute an instruction —another ALU is performing— or stay **idle**.

4.6.1 AVX registers

Listing 4.1: Intrinsics

```
//AVX-Programming with C/C++ Intrinsics
__m256 a, b, c; // declare AVX registers
... // initialize a and b
c = _mm256_add_ps(a, b); // c[0:8] = a[0:8] + b[0:8]
// or c = _mm512_add_pd(a, b); c[0:8] = a[0:8] + b[0:8]
```

Intrinsics are assembly-coded functions that can be used in C/C++ to exploit SIMD parallelism. They provide a light abstraction from assembly. Exploiting intrinsics as in Lst. 4.1 may lead to a 8x speedup, they are very powerful.

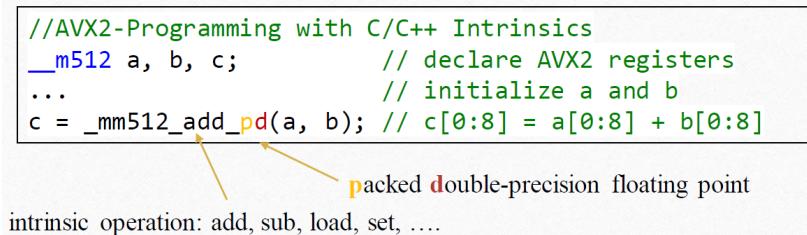


Figure 4.6: Intrinsics explained

There are further notions on intrinsics and vectorization I did not report here

4.7 Vectorizing code

4.7.1 AoS vs SoA



Figure 4.7: AoS vs SoA

- ◇ **AoS - Array of Structures:** stores records consecutively in a single array.
Copilot: “each element of the array is a struct, and the struct contains all the data.”
- ◇ **SoA - Structure of Arrays:** uses one array per dimension. Each array stores only the values of the associated element dimension
Copilot: “each element of the array is a single data type, and the data is stored in different arrays.”

Typically AoS does not allow for vectorization, since the data is not contiguous in memory, while SoA does, since the data is contiguous, allowing to perform operations on up to #registers triplets of data.

4.7.2 Loops - Compiler Optimization

The compiler is able to optimize the code and vectorize it, but it is not always possible. In particular, there are some guidelines concerning loops to allow the compiler to vectorize the code:

- ◇ The loop count must be known at entry to the loop.
- ◇ Conversely, the exit condition must not be data-dependent.
 - **break, goto, switch, return** statements are forbidden
 - **if** is allowed and may be vectorized if it can be implemented as a masked assignment
- ◇ No library functions calls. If a function call may be inlined it is OK.
What does this mean?
- ◇ No read/write dependencies between iterations. e.g. the following code is *not* vectorizable:


```
|     A[0] = 0; for(int_i=1; i<N; ++i) A[i]= A[i-1] + 1;
```
- ◇ No pointer aliasing, i.e. no two pointers can point to the same memory location.
 - The type qualifier **_restrict_** tells the compiler that the pointer is the only pointer that points to that memory location.

Chapter 5

Distributed Memory Systems

The first issue to address is the communication between different nodes. The main problem is that the memory is distributed, so we need to find a way to communicate between different nodes.

5.1 Interconnection Networks

Interconnection networks for parallel systems share many technical features of WAN, but they have very different requirements:

- ◊ $\mathcal{O}(10^1 \div 10^5)$ nodes
- ◊ Distances ranging in $\mathcal{O}(10^0 \div 10^1)$ meters

Key metrics for us are:

- ◊ **Latency** - time lapse between the instant a packet starts to be transmitted and the instant it is —entirely— received
 - **No-load** (or *Zero-load*) latency: the latency experienced by a packet when there is no other traffic on the network
 - **Under-load** latency: the latency experienced by a packet when there is other traffic on the network, but below the *saturation point*
Saturation point is a crucial threshold we will discuss in future.
- ◊ (Offered) **Throughput** - the amount of data that transmitted over the network in a given time. The **Saturation Throughput** is the maximum amount of traffic sustained by the network, it is the point at which the network is fully loaded.
- ◊ **Bandwidth** - theoretical maximum data transfer rate under ideal conditions on a network path.

Two aspects which are fundamental but which we won't address are:

- ◊ **Routing**
- ◊ **Flow control**

5.1.1 Terminology

Endpoints, links, switches...

- ◊ **Direct Network** aka *static* - nodes are both endpoints and switches
- ◊ **Indirect Network** aka *dynamic* - endpoints are connected indirectly through switches
- ◊ **Degree** - number of maximum neighbors a node can have
- ◊ **Diameter** - longest of shortest paths between any two nodes
maximum (minimum) distance between two nodes
- ◊ **Bisection width** - minimum number of links that must be removed to split the network into two equal halves

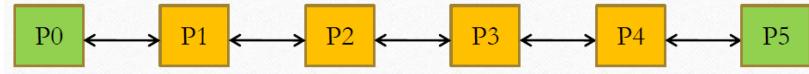


Figure 5.1: Linear Array topology
 $\text{degree} = 2$ $\text{Diameter} = n - 1$ $\text{bw} = 1$

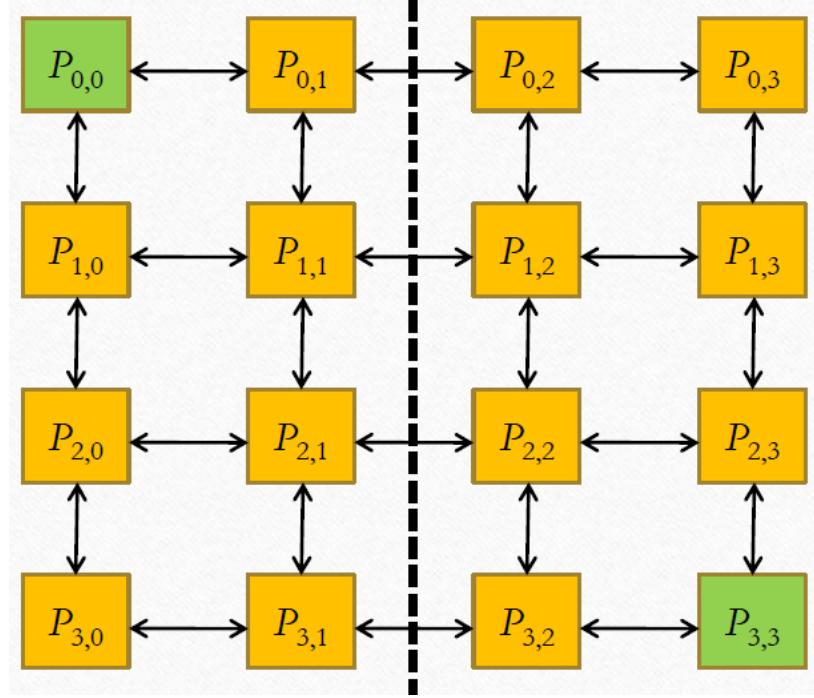


Figure 5.2: 2D Mesh topology
 $M(d, d)$ has $n = d^2$ endpoints $\text{degree} = 4$ $\text{Diameter} = 2(d - 1) = \mathcal{O}(\sqrt{n})$ $\text{bw} = d = \mathcal{O}(\sqrt{n})$

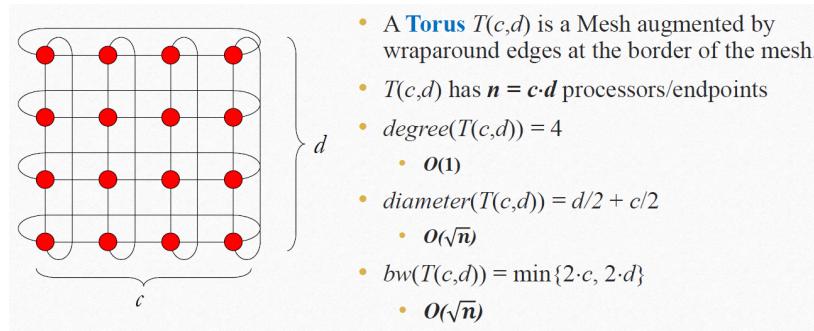


Figure 5.3: Torus topology

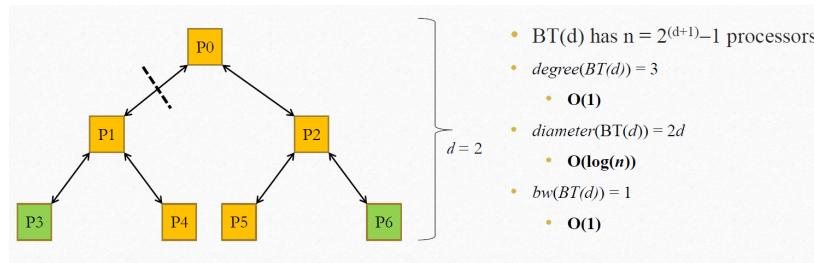


Figure 5.4: Binary Tree topology

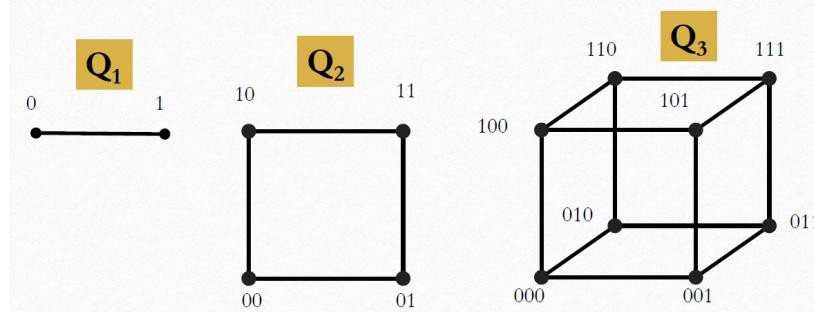


Figure 5.5: Hypercube topology

5.1.2 Examples

5.1.2.1 Hypercube

The hypercube $\mathcal{Q}_d (d \geq 1)$ is the graph that has vertices representing the 2^d binary strings of length d . Two vertices are adjacent if their binary strings differ in *exactly* one bit.

The hypercube is a very interesting topology because it has a very low diameter and bisection width, but it is very expensive to build.

5.1.3 Criteria for Evaluation - Summary

So, what do we want from an interconnection network?

- ◊ **Constant degree** - Allows scaling the network without increasing the degree of the nodes
- ◊ **Low Diameter** - Allows for low latency
- ◊ **High bisection width** - Allows for high bandwidth, but may require to have dynamic degree.

5.1.3.1 Simple Communication cost model

$$T_{comm} = t_0 + n \times s \approx \begin{cases} t_0 & \text{for small } n \\ n \times s & \text{for large } n \end{cases} \quad (5.1)$$

- ◊ t_0 start-up time (network and communication runtime setup)
 - Includes all costs for sending the shortest message. (*latency*)
 - Its value may be different for different programming models
- ◊ n amount of (e.g. bytes) data to be transferred
- ◊ s transmission cost. Usually $s = \frac{1}{B}$ with B available bandwidth on the path
 - s is limited by the slowest part of the path between the sender and the receiver processes
 - s includes both SW agents (e.g. app max output rate) and HW agents (e.g. network max throughput)

5.1.4 More advanced topologies

- ◊ **Fat-tree** - a tree-like topology that has a high bisection width
- ◊ **Dragonfly** - a topology that combines a tree-like structure with a mesh-like structure
- ◊ **Jellyfish** - a topology that is based on a random graph

5.1.4.1 Fat-tree

All leaves are endpoints and all internal nodes are switches. Links higher in the tree have higher bandwidth. The idea is to keep the bandwidth constant across the tree.

The key issue is the cost, which increases with the depth of the tree. To overcome this problem, we can use a **Spine-Leaf** topology, where the tree is split into two parts: the spine and the leaves.

In the slides this is considered as a different implementation of the fat-tree, but it is actually a different topology.

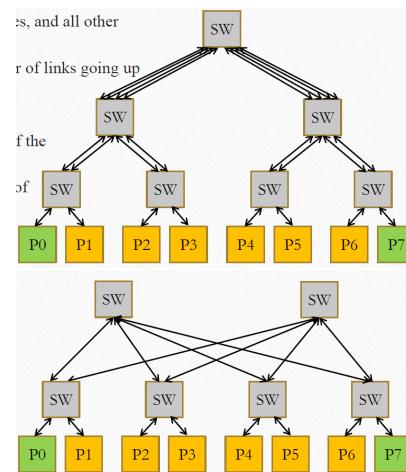


Figure 5.6: Fat tree and Spine Leaf

5.1.4.2 Dragonfly

This is the most used network used in supercomputers. It is divided in three levels:

1. Router
2. Group
3. System

Each router has connections to p endpoints, $a - 1$ local channels (to other routers in the same group) and h global channels (to routers in other groups).

A group consists of a routers and group has ap connections to endpoints, and ah connections to other groups.

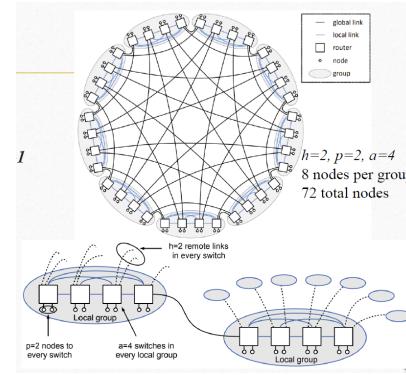


Figure 5.7: Dragonfly

5.2 Computation-to-Communication Overlap

It is always useful for the sender to execute a **non-blocking send** to the destination process: the NICE executes the data transfer in the background, and then notifies the sender process when the transmission is completed, so that the sender can continue with its computation. This creates a partial or full *overlap* between computation and communication.

How can we achieve this overlap?

5.2.1 Synchronous vs Asynchronous Communication

The **asynchrony degree** of a channel is the maximum number of messages that can be sent on the channel without waiting for the receiver to start receiving data. Essentially, it is the length of the queue of messages that can be sent on the channel.

- ◊ **Synchronous** - the operation completes only after the message has been received/sent. i.e. a peer **blocks** until the other peer completes the operation.
- ◊ **Asynchronous** - the communication operation returns immediately, without waiting for the message to be effectively sent/received.
 - The completion/success will be tested later on
 - The number of messages that can be sent depends on the **asynchrony degree** of the channel.

Note also that sometimes **input non-determinism** must be taken into account: the *receive* operation may return a message from any channels in the set, and the message picking algorithm is random and non-predictable.

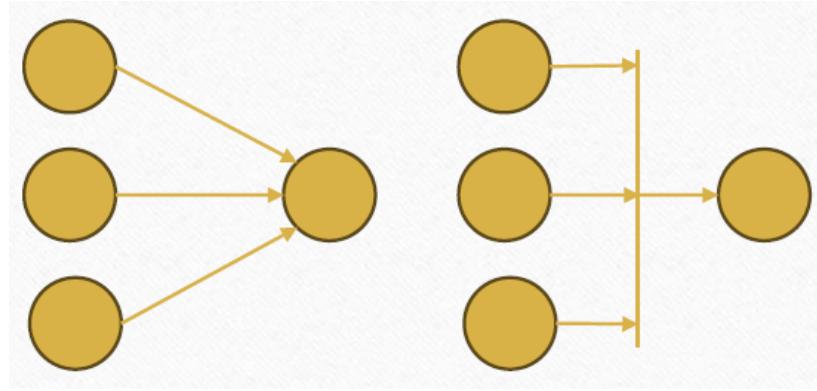


Figure 5.6: Input non-determinism

5.3 Foster's Parallel Algorithm Design Method

There is no one-fits-all recipe for parallelizing sequential programs, however... Foster's **PCAM** approach provides some useful guidelines to keep in mind when designing parallel algorithms.

- ◊ **Partitioning** - divide the problem into a large amount of small fine-grained tasks that can be executed in parallel
- ◊ **Communication** - determine the required communication pattern between the tasks (*data dependencies*)
- ◊ **Agglomeration** - combine identified tasks into larger coarse-grained tasks to reduce communication by improving data locality
- ◊ **Mapping** - assign tasks to processes and threads to minimize communication, enable concurrency and *balance workload*